



Project C places conflation



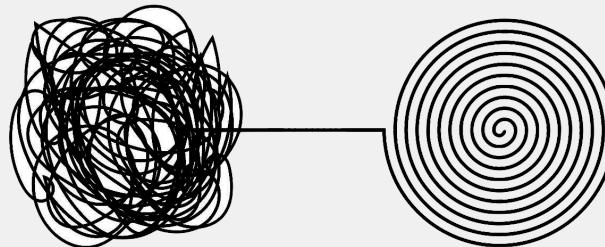
Problem Statement



Determine if small, scalable Language Models can perform better than the current places classifier.

1. Improve Accuracy: Can a model achieve higher accuracy in matching places data?
2. Improve Efficiency: Can a model offer a better price-to-performance ratio than the existing matcher?

Problem Breakdown



Analyzing Cost:

Time it takes to run – how long the full evaluation requires

Tokens used per match – how many LLM tokens are consumed for each comparison

Cost per match – approximate monetary cost for processing each pair

Analyzing Performance:

1. **Accuracy** – how many predictions were correct overall
2. **Precision** – of the predicted matches, how many were truly matches
3. **Recall** – of the true matches, how many were found by the model
4. **F1 Score** – balanced measure combining precision and recall

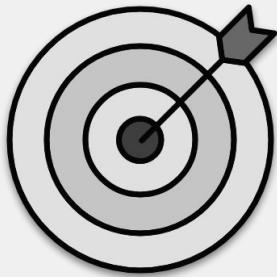
Constraints + Challenges

Must work at large scale – system needs to process thousands of POIs

Must stay within budget – model choice and token usage must remain cost-effective

Place data is noisy – spelling, formatting, and abbreviations vary widely

Global variation makes matching harder – addresses differ across countries and languages



Objective:

Replace Overture's existing place matcher with a scalable, language-model based system that improves matching precision and cost efficiency.

01

Key Result

Train and evaluate at least three scalable models on the dataset given.

02

Key Result

Reduce average cost per inference by $\geq 15\%$, as measured using compute time and API/token usage logs.

03

Key Result

Reduce false matches by $\geq 10\%$ relative to the overture matching model's baseline metrics.

OKR

Completion

KR1 – Passed

Evaluated four different scalable models, exceeding the requirement of three.
Models were tested using a consistent validation pipeline, allowing direct comparison.

KR2 – Passed

Achieved a major reduction in cost per inference.
The best-performing model (llama-3.1-8b-instant) costs significantly less than many other larger models.
Strong performance despite dramatically lower cost makes this KR fully met.

KR3 – Almost Met

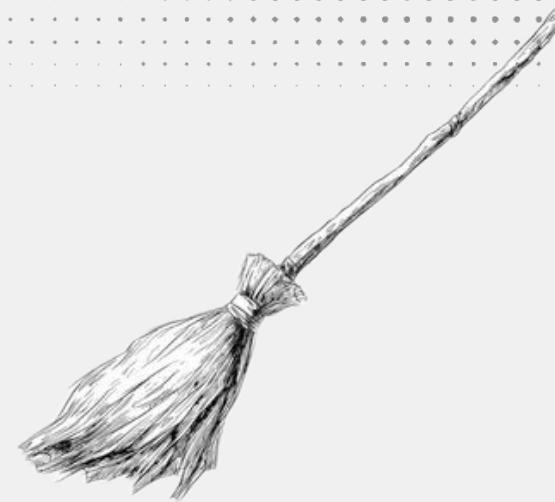
None of the tested models reached or exceeded Overture's baseline precision of 0.93.
Primary limitation: smaller and mid-size LLMs struggle with borderline cases and subtle address distinctions, which Overture's established model handles better.
With further tuning, prompting, or ensemble methods, this KR could become reachable.

Approach

Data Cleaning	Similarity Score	LLM Matching	Evaluate Results
01 Data cleaning ensures all fields follow a consistent format so the matching models can compare records reliably. It reduces noise, fixes inconsistencies, and prepares the data for evaluation.	02 Similarity scoring provides an initial estimate of how similar two POI records are. It serves as a fast baseline classifier before the LLM performs deeper reasoning.	03 Use a language model to validate or correct the fuzzy match by checking the key address fields. It outputs a JSON decision with a same_place label, confidence score, and a brief explanation.	04 Evaluate LLM predictions using accuracy, precision, and recall, then review incorrect matches to understand issues.

Data Cleaning

- Expand abbreviations
 - st = street
 - ave = avenue
 - cir = circle
 - N = north
 - apt = apartment
- convert to lowercase
- remove extra characters from phone number (such as + and ())
- remove the https: from websites
- Example normalized data:
`{"freeform":"Avenida Carlos Cano","locality":"Los Barrios","postcode":"11370","country":"ES"}`
- **avenida carlos cano,los barrios,,11370,ES**



Fuzzy Match Score

- Calculate a fuzzy matching similarity score for each field in each row
- RapidFuzz token_sort_ratio
 - Compares words ignoring order (e.g., "Starbucks Coffee" ≈ "Coffee Starbucks").
- RapidFuzz partial_ratio
 - Handles substring overlap (e.g., "McDonalds" ≈ "McDonalds Hamburgers").
- RapidFuzz token_set_ratio
 - Handles cases with repeated/extraneous tokens (e.g., "CVS Pharmacy" ≈ "CVS").
- Jaro-Winkler similarity
 - Character-level similarity robust to typos (e.g., "Walgreens" ≈ "Walgreans").
- Sum up scores for each field for a total similarity score

similarity =
20% name_sim
50% address_sim
15% phone_sim
5% website_sim
10% category_sim



Base Accuracy

- Using just a fuzzy similarity score, and choosing 1 or 0 based on a threshold of 0.671 we get a base accuracy = ~74%

```
Cleaning data...
```

	address_freeform	base_address_freeform	match_score	label	pred_label
0	avenida carlos cano	avenue carlos cano glorieta la montera	0.655556	1.0	1
1	609 west maple avenue	609 west maple avenue	0.723116	0.0	1
2	688 9th avenue	445 west 45th street	0.382655	0.0	0
3	401 east chestnut street unit 790	401 east chestnut street unit 510	0.671309	0.0	1
4	avenida capitao inacio prata sn	rua capitao inacio prata south north	0.678917	1.0	1

```
Accuracy: 0.7403333333333333  
(crwn102) eduroam-169-233-247-85:pipeline indiram$ 
```

Prompt Creation

1. PRIORITIZE "sim" AND "pred":

- If pred = 1 and sim is high (≥ 0.7), assume same_place = 1 unless the address clearly disagrees.
- If pred = 0 and sim is low (≤ 0.3), assume same_place = 0 unless the address clearly shows they match.

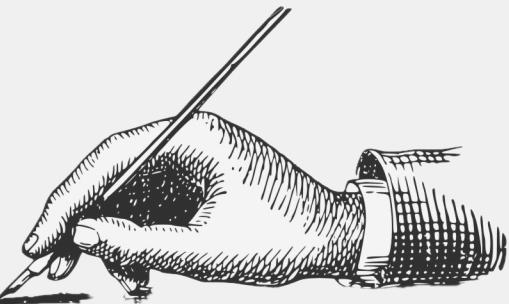
2. USE THE ADDRESS AS A CONFIRMATION CHECK (SECONDARY):

- Address fields are already normalized.
- If the baseline decision from sim + pred conflicts with the address (e.g., different street numbers or different cities), flip the decision.

3. GENERAL RULE:

- Similar address → supports same_place = 1
- Very different address → supports same_place = 0
- When uncertain, rely on sim + pred unless the address strongly contradicts it.





Prompt Cont.

OUTPUT REQUIREMENTS:

- Return ONLY a JSON object.
- "same_place" must be 1 or 0.
- "confidence" must be a number between 0.0 and 1.0.
- "explanation" must be 1 sentence and MUST describe the specific fields that led to the decision.
- Do NOT reuse generic phrases across answers.

RECORD A:

Address (freeform): {address_freeform}

Locality / City: {address_locality}

Region / State: {address_region}

Postcode: {address_postcode}

Country: {address_country}

Phone: {phone}

RECORD B:

.....

Outputs

llama-3.1-8b-instant

- The inputs are fed into the model, and we are given an output in the format we specified
- Explanation is added for further context on why the model made its decision
- can help finetune the prompt to account for mistakes

Example 1

avenida carlos cano,los barrios,
avenue carlos cano glorieta la montera,los barrios

```
{'same_place': 1, 'confidence': 0.9, 'explanation': 'The  
decision to consider these records as the same place  
was made because the street name, locality, postcode,  
and country matched, and although the phone numbers  
were not identical, the similarity score and the  
matching address components supported this  
conclusion.'}
```

Example 2

1100 west stewart drive,orange,ca
1000 west la veta drive 1st floor,orange,ca

```
{'same_place': 0, 'confidence': 0.6, 'explanation': "The  
decision was made because the street numbers and  
street names in the address fields of the two records  
do not match, with record A having '1100 west stewart  
drive' and record B having '1000 west la veta drive 1st  
floor', which indicates they are likely different places."}
```

Outputs...

Translation Issue: Guessed 0, Actual: 1

avenida **capitao inacio prata sn**,sao goncalo do amarante,ce,62670000,BR,,

rua **capitao inacio prata south north**,,62670000,BR,8533156140.0,0.7104214811
4427

Reasoning: "The decision was based on the conflicting street types and names, with 'avenida' in Record A and 'rua' in Record B, as well as the missing locality and region in Record B, which outweighed the matching postcode and country."

Address Overlap Issue: Guessed: 1, Actual: 0

apollo children hospital no 15 shafee mohammed road thousand lights,chennai,tamil nadu,600006,IN,8605001066.0,

apollo hospitals greams road no 21 greams lane off greams road,chennai,tamil nadu,
600006,IN,8605001066.0,0.5454377901122027,1,0.8,

Reasoning: "The decision to consider these records as the same place was made because the phone number, locality, postcode, and country matched exactly between the two records, despite some differences in the address freeform field."

Results

Model Token Costs

1. llama-3.1-8b-instant – **\$0.31** per 1M
2. llama-4-scout-17b – **\$0.85** per 1M
3. llama-3.3-70b – **\$5.17** per 1M
4. kimi-instruct – **\$9.45** per 1M
5. kimi-0905 – **\$13.42** per 1M

Model Speeds:

1. kimi-k2-instruct-0905 – 0.222s avg (4.50 matches/sec)
2. kimi-k2-instruct – 0.264s avg (3.78 matches/sec)
3. llama-3.1-8b-instant – 0.413s avg (2.42 matches/sec)
4. llama-4-scout-17b-16e-instruct – 1.083s avg (0.92 matches/sec)



Accuracy



77%

llama-3.1-8b-instant
Accuracy: 0.77
Precision: 0.7651
Recall: **0.912**
F1 Score: 0.8321



79%

kimi-k2-instruct-0905
Accuracy: 0.795
Precision: 0.8182
Recall: 0.864
F1 Score: 0.8405



78%

llama-4-scout-17b-16e
-instruct
Accuracy: 0.785
Precision: **0.8534**
Recall: 0.792
F1 Score: 0.8216



80%

kimi-k2-instruct
Accuracy: **0.800**
Precision: 0.8195
Recall: 0.872
F1 Score: **0.845**



llama-3.1-8b-instant

1. Highest recall (0.912)
finds the most true matches, minimizes missed duplicates.
2. Strong F1 score (0.8321)
good balance of precision and recall.
3. Lowest cost (\$0.31 per 1M tokens)
much cheaper than other tested models
4. Fast runtime (0.413s avg)
efficient for large-scale processing, scalable
5. Stable precision (0.7651)
good reliability without excessive false positives

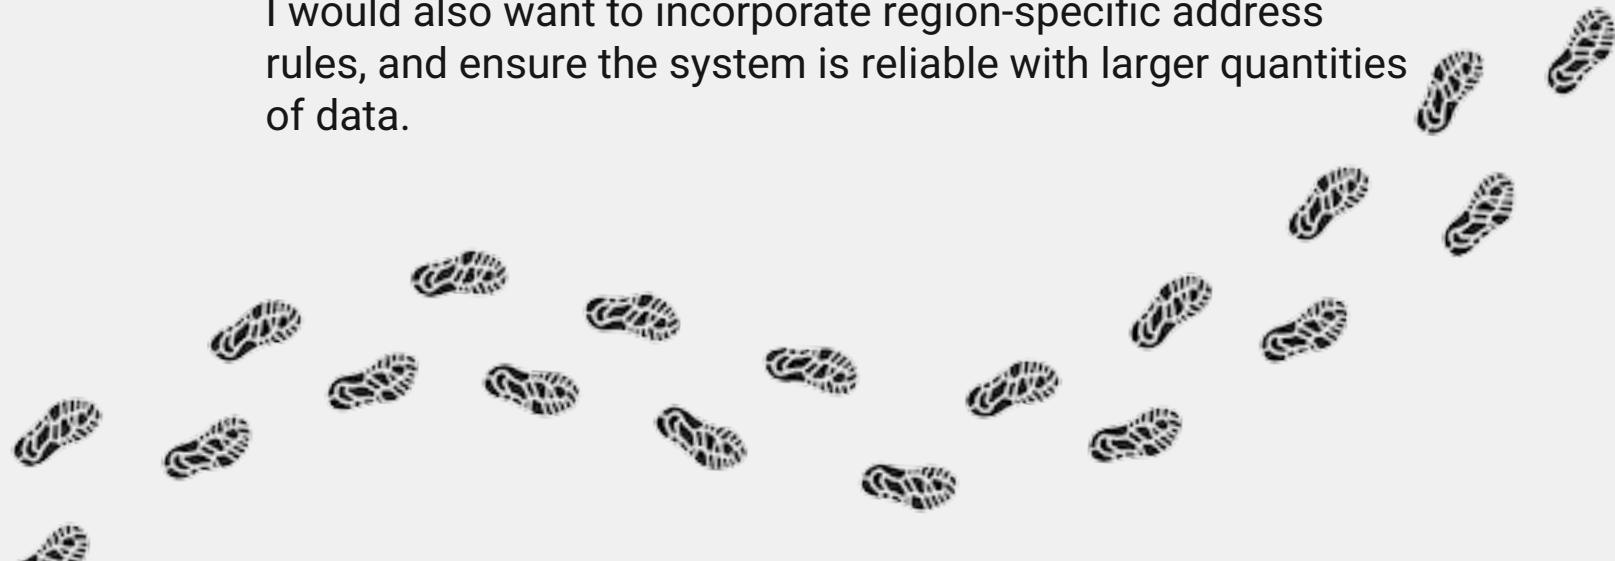
Overall Best

Next Steps

To continue this project, I want to tune the parameters even more and refine the model prompts to reduce common errors.

I also plan to experiment with combining multiple model outputs maybe using a framework like XGBoost to create a ensemble classifier.

I would also want to incorporate region-specific address rules, and ensure the system is reliable with larger quantities of data.





Thank you

Do you have any questions?

[Github Repo](#)